

✓ Importing Necessary Libraries

```
import random
import numpy as np
from tensorflow.keras.datasets import cifar10
import matplotlib.pyplot as plt
from sklearn.decomposition import PCA
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, classification_report

import warnings
warnings.filterwarnings('ignore')
```

✓ Loading the dataset

```
import tensorflow as tf
from keras.datasets import cifar10
from tensorflow.keras.utils import to_categorical

# Load the CIFAR-10 dataset
(train_images, train_labels), (test_images, test_labels) = cifar10.load_data()

# Preprocess the data
train_images = train_images.astype('float32') / 255.0
test_images = test_images.astype('float32') / 255.0
train_labels = to_categorical(train_labels)
test_labels = to_categorical(test_labels)

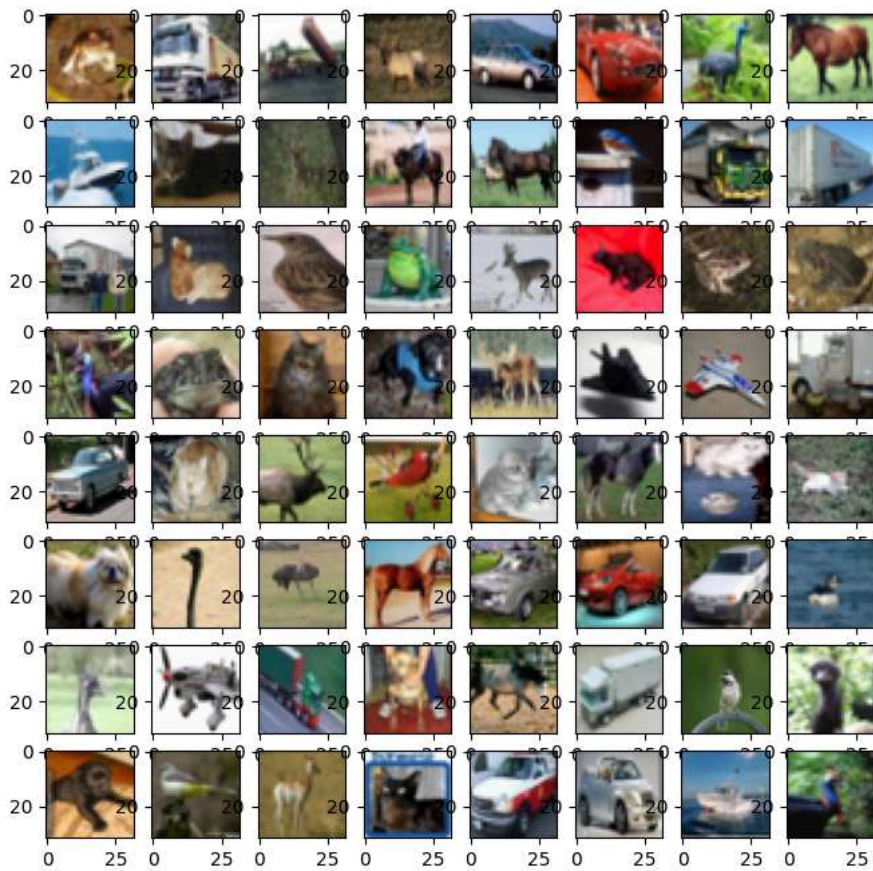
Downloading data from https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz
170498071/170498071 [=====] - 4s 0us/step

# Checking the size of the test dataset
test_images.shape

(10000, 32, 32, 3)
```


✓ Checking some of the images loaded in our dataset

```
fig = plt.figure(figsize = (8,8))
for i in range(64):
    ax = fig.add_subplot(8,8,i+1)
    ax.imshow(train_images[i], cmap=plt.cm.bone)
plt.show()
```



```
# example of loading the cifar10 dataset
from matplotlib import pyplot
from keras.datasets import cifar10
# load dataset
(trainX, trainy), (testX, testy) = cifar10.load_data()
# summarize loaded dataset
print('Train: X=%s, y=%s' % (trainX.shape, trainy.shape))
print('Test: X=%s, y=%s' % (testX.shape, testy.shape))
# plot first few images
for i in range(9):
    # define subplot
    pyplot.subplot(330 + 1 + i)
    # plot raw pixel data
    pyplot.imshow(trainX[i])
# show the figure
pyplot.show()
```

Train: X=(50000, 32, 32, 3), y=(50000, 1)
 Test: X=(10000, 32, 32, 3), y=(10000, 1)



- ✓ Our Image initially having (50000, 32x32x3) is reshaped to (50000 x 3072) dimensions, for further analysis

```
x_train = trainX.reshape(train_images.shape[0],-1)
x_test = testX.reshape(test_images.shape[0], -1)
```

```
x_train.shape

(50000, 3072)
```

- ✓ We use Robust Scaler to scale the dataset. Robust Scaler is considered to be more robust to outliers

```
from sklearn.preprocessing import StandardScaler, RobustScaler
sc = RobustScaler()
x_train_scaled = sc.fit_transform(x_train)
```

```
x_test_scaled = sc.transform(x_test)
```

```
y_train = trainy
y_test = testy
```

```
# pca = PCA()
# pca.fit_transform(x_train_scaled)
```

```
# k = 0
# total = sum(pca.explained_variance_)
# current_sum = 0

# while(current_sum / total < 0.99):
#     current_sum += pca.explained_variance_[k]
#     k += 1
# k
```

```
# pca_cifar = PCA(n_components=k,whiten=True)
# x_train_pca_cifar = pca_cifar.fit_transform(x_train_scaled)
# x_test_pca_cifar = pca_cifar.transform(x_test_scaled)
```

```
# x_test_pca_cifar.shape
```

```
# y_train.shape
```

✓ Let us now fit model and check for evaluation metrics

```
from sklearn.linear_model import LogisticRegression
lr = LogisticRegression()
lr.fit(x_train_scaled, y_train.ravel())
```

```
▼ LogisticRegression
LogisticRegression()
```

```
y_pred_lr = lr.predict(x_test_scaled)
logistic_regression_score = accuracy_score(y_test, y_pred_lr)
logistic_regression_score

0.4014
```

```
from sklearn.metrics import precision_score, recall_score, f1_score, classification_report
```

```
y_pred_lr = lr.predict(x_test_scaled)

precision = precision_score(y_test, y_pred_lr, average = 'weighted')
recall = recall_score(y_test, y_pred_lr, average = 'weighted')
f1 = f1_score(y_test, y_pred_lr, average = 'weighted')
```

```
classification_rep = classification_report(y_test, y_pred_lr)
```

```
print("Precision:", precision)
print("Recall:", recall)
print("F1 Score:", f1)
print("Classification Report:\n", classification_rep)
```

```
Precision: 0.3985215429260664
Recall: 0.4014
F1 Score: 0.3991216365736525
Classification Report:
      precision    recall  f1-score   support

0         0.46         0.48         0.47         1000
1         0.47         0.48         0.47         1000
2         0.32         0.28         0.30         1000
3         0.28         0.27         0.27         1000
4         0.36         0.30         0.33         1000
5         0.33         0.34         0.33         1000
6         0.39         0.45         0.42         1000
7         0.45         0.43         0.44         1000
8         0.49         0.53         0.51         1000
9         0.44         0.46         0.45         1000

 accuracy          0.40         10000
 macro avg         0.40         0.40         0.40         10000
 weighted avg         0.40         0.40         0.40         10000
```

```
from sklearn.metrics import confusion_matrix
```

```
conf_matrix = confusion_matrix(y_test, y_pred_lr)
```

```
print("Confusion Matrix:")
print(conf_matrix)
```

```
Confusion Matrix:
[[484  41  53  41  20  27  25  51 182  76]
 [ 52 478  31  39  28  40  41  44  78 169]
 [ 97  39 280 105 119  82 131  76  47  24]
 [ 38  50  93 267  61 195 144  54  37  61]
 [ 56  24 137  69 297  92 168 105  26  26]
 [ 40  46  96 172  88 337  89  68  36  28]
 [ 22  41  75 144  91  97 453  36  16  25]
 [ 53  48  66  71  94  87  43 428  31  79]
 [154  72  28  28  12  38  18  22 529  99]
 [ 65 178  19  30  17  27  47  60  96 461]]
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
rf_classifier = RandomForestClassifier(n_estimators = 50)
```

```
rf_classifier.fit(x_train_scaled, y_train.ravel())
```

```
▼ RandomForestClassifier
RandomForestClassifier(n_estimators=50)
```

```
y_pred_rf = rf_classifier.predict(x_test_scaled)
```

```
rf_classifier_score = accuracy_score(y_test, y_pred_rf)
```

```
rf_classifier_score
```

```
0.4508
```

```
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, classification_report, confusion_matrix
```

```
y_pred_rf = rf_classifier.predict(x_test_scaled)
```

```
accuracy_rf = accuracy_score(y_test, y_pred_rf)
```

```
precision_rf = precision_score(y_test, y_pred_rf, average = 'weighted')
```

```
recall_rf = recall_score(y_test, y_pred_rf, average = 'weighted')
```

```
f1_rf = f1_score(y_test, y_pred_rf, average = 'weighted')
```

```
classification_rep_rf = classification_report(y_test, y_pred_rf)
```

```
conf_matrix_rf = confusion_matrix(y_test, y_pred_rf)
```

```
print("Random Forest Classifier Metrics:")
```

```
print("Accuracy:", accuracy_rf)
```

```
print("Precision:", precision_rf)
```

```
print("Recall:", recall_rf)
```

```
print("F1 Score:", f1_rf)
```

```
print("Classification Report:\n", classification_rep_rf)
```

```
print("Confusion Matrix:")
```

```
print(conf_matrix_rf)
```



```
Random Forest Classifier Metrics:
```

```
Accuracy: 0.4508
```

```
Precision: 0.4465691128922581
```

```
Recall: 0.4508
```

```
F1 Score: 0.44759177397875494
```

```
Classification Report:
```

	precision	recall	f1-score	support
0	0.52	0.57	0.54	1000
1	0.50	0.54	0.52	1000
2	0.33	0.32	0.33	1000
3	0.31	0.25	0.28	1000
4	0.38	0.38	0.38	1000
5	0.42	0.38	0.40	1000
6	0.47	0.53	0.50	1000
7	0.48	0.42	0.45	1000
8	0.59	0.59	0.59	1000
9	0.47	0.52	0.49	1000
accuracy			0.45	10000
macro avg	0.45	0.45	0.45	10000
weighted avg	0.45	0.45	0.45	10000

```
Confusion Matrix:
```

```
[[566 39 63 24 30 20 21 24 156 57]
 [ 43 542 24 38 28 31 30 34 54 176]
 [116 37 320 76 163 64 107 58 26 33]
 [ 59 40 90 254 83 185 135 71 27 56]
 [ 45 30 162 52 384 50 142 83 29 23]
 [ 31 39 91 168 80 380 80 78 23 30]
 [ 18 35 107 72 107 57 530 34 5 35]
 [ 47 45 71 71 105 78 48 422 19 94]
 [113 87 20 33 20 23 11 23 592 78]
 [ 54 185 15 36 20 16 33 44 79 518]]
```

▼ This time we check for a validation set of 20-80. We do this for one model i.e, Logistic Regression

```

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

# Split the data into training and testing sets
x_val_train, x_val_test, y_val_train, y_val_test = train_test_split(x_train, y_train, test_size=0.2, random_state=42)

# Fit the model
logreg = LogisticRegression()
logreg.fit(x_val_train, y_val_train)

# Make predictions
y_val_pred = logreg.predict(x_val_test)

# Calculate evaluation metrics
accuracy_val = accuracy_score(y_val_test, y_val_pred)
precision_val = precision_score(y_val_test, y_val_pred, average = 'macro')
recall_val = recall_score(y_val_test, y_val_pred, average = 'macro')
f1_val = f1_score(y_val_test, y_val_pred, average = 'macro')

print("Accuracy:", accuracy_val)
print("Precision:", precision_val)
print("Recall:", recall_val)
print("F1 Score:", f1_val)

    Accuracy: 0.4045
    Precision: 0.3990480934172669
    Recall: 0.40461833231965516
    F1 Score: 0.4001986792394578

```

Applying Cross - validation with different k values. Choosing only logistic regression model for comparison to reduce computational intensity

✓ For k=3

```

from sklearn.model_selection import cross_val_score, KFold
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

# Create a logistic regression model
logreg = LogisticRegression()

# Perform 5-fold cross-validation
cv = KFold(n_splits=3, shuffle=True, random_state=42)

# Calculate evaluation metrics using cross-validation
accuracy_val = cross_val_score(logreg, x_train, y_train, cv=cv, scoring='accuracy').mean()
# precision_val = cross_val_score(logreg, x_train, y_train, cv=cv, scoring='precision_macro').mean()
# recall_val = cross_val_score(logreg, x_train, y_train, cv=cv, scoring='recall_macro').mean()
# f1_val = cross_val_score(logreg, x_train, y_train, cv=cv, scoring='f1_macro').mean()

print("Accuracy:", accuracy_val)
print("Precision:", precision_val)
print("Recall:", recall_val)
print("F1 Score:", f1_val)

    Accuracy: 0.40174007392287053
    Precision: 0.3990480934172669
    Recall: 0.40461833231965516
    F1 Score: 0.4001986792394578

```

✓ For k=5

```

from sklearn.model_selection import cross_val_score, KFold
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

# Create a logistic regression model
logreg = LogisticRegression()

# Perform 5-fold cross-validation
cv = KFold(n_splits=5, shuffle=True, random_state=42)

# Calculate evaluation metrics using cross-validation
accuracy_val = cross_val_score(logreg, x_train, y_train, cv=cv, scoring='accuracy').mean()
# precision_val = cross_val_score(logreg, x_train, y_train, cv=cv, scoring='precision_macro').mean()
# recall_val = cross_val_score(logreg, x_train, y_train, cv=cv, scoring='recall_macro').mean()
# f1_val = cross_val_score(logreg, x_train, y_train, cv=cv, scoring='f1_macro').mean()

print("Accuracy:", accuracy_val)
# print("Precision:", precision_val)
# print("Recall:", recall_val)
# print("F1 Score:", f1_val)

Accuracy: 0.40546
Precision: 0.3990480934172669
Recall: 0.40461833231965516
F1 Score: 0.4001986792394578

```

✓ For k = 10

```

from sklearn.model_selection import cross_val_score, KFold
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

# Create a logistic regression model
logreg = LogisticRegression()

# Perform 5-fold cross-validation
cv = KFold(n_splits=10, shuffle=True, random_state=42)

# Calculate evaluation metrics using cross-validation
accuracy_val = cross_val_score(logreg, x_train, y_train, cv=cv, scoring='accuracy').mean()
# precision_val = cross_val_score(logreg, x_train, y_train, cv=cv, scoring='precision_macro').mean()
# recall_val = cross_val_score(logreg, x_train, y_train, cv=cv, scoring='recall_macro').mean()
# f1_val = cross_val_score(logreg, x_train, y_train, cv=cv, scoring='f1_macro').mean()

print("Accuracy:", accuracy_val)
# print("Precision:", precision_val)
# print("Recall:", recall_val)
# print("F1 Score:", f1_val)

Accuracy: 0.40574000000000005
Precision: 0.3990480934172669
Recall: 0.40461833231965516
F1 Score: 0.4001986792394578

```

✓ As a part 2 of our problem, we repeat the same experiment on grayscale images this time.

```

import numpy as np
import tensorflow as tf

# Load CIFAR-10 dataset
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.cifar10.load_data()

# Convert images to grayscale
x_train_gray = np.mean(x_train, axis=3)
x_test_gray = np.mean(x_test, axis=3)

# Add a single channel dimension
x_train_gray = np.reshape(x_train_gray, x_train_gray.shape + (1,))
x_test_gray = np.reshape(x_test_gray, x_test_gray.shape + (1,))

# Normalize pixel values to the range [0, 1]
x_train_gray = x_train_gray.astype('float32') / 255.0
x_test_gray = x_test_gray.astype('float32') / 255.0

# Flatten the images
x_train_flat = x_train_gray.reshape(x_train_gray.shape[0], -1)

from sklearn.linear_model import LogisticRegression
lr = LogisticRegression()
lr.fit(x_train, y_train.ravel())

    ▾ LogisticRegression
    LogisticRegression()

from sklearn.metrics import precision_score, recall_score, f1_score, classification_report

y_pred_lr = lr.predict(x_test)

precision = precision_score(y_test, y_pred_lr, average = 'weighted')
recall = recall_score(y_test, y_pred_lr, average = 'weighted')
f1 = f1_score(y_test, y_pred_lr, average = 'weighted')

classification_rep = classification_report(y_test, y_pred_lr)

print("Precision:", precision)
print("Recall:", recall)
print("F1 Score:", f1)
print("Classification Report:\n", classification_rep)

Precision: 0.3057978481452976
Recall: 0.3124
F1 Score: 0.3077123775021974
Classification Report:
      precision    recall  f1-score   support

0         0.34        0.36        0.35         1000
1         0.37        0.39        0.38         1000
2         0.24        0.20        0.22         1000
3         0.22        0.16        0.19         1000
4         0.25        0.22        0.24         1000
5         0.32        0.32        0.32         1000
6         0.29        0.31        0.30         1000
7         0.30        0.30        0.30         1000
8         0.34        0.40        0.37         1000
9         0.38        0.45        0.42         1000

 accuracy          0.31         10000
 macro avg          0.31         10000
 weighted avg       0.31         10000

rf = RandomForestClassifier()
rf.fit(x_train, y_train.ravel())

    ▾ RandomForestClassifier
    RandomForestClassifier()

```